2_sim_rl.py

1. <u>For all bots which are currently at a node & idle</u>

Update:

- true idleness of all nodes in graph (true) = +Δt
- expected idleness of all edges in the graph (expect) = +Δt

<u>Calculate</u>: here, learning rate (α)= 0.1, discount factor (γ)= 0.95

- Value function all edges where bots are present (Q) =

$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \overbrace{\Big(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{old value}} \Big)}^{\text{temporal difference}}$$
$$\underbrace{\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad}_{\text{new value (temporal difference target)}}$$

- Reward (r_t)= $log(\,|expect \;-\; true|\,)$and

  $r_t = 0$ if [expect = true]

- Softmax of Value function = value_exp = $\frac{e^{Q_i}}{\sum_{j=i}^{k} e^{Q_j}}$ (summation over all edges)

Set:

- True idleness of nodes where bots are present = 0
- Expected idleness of nodes wrt the corresponding bot = 0

NOTE: So effectively, we are calculating expected idleness perceived by the bot as the time elapsed since this bot last visited the node.
This is the **AGENT model** where the bot is calculating expected idleness by seeing the time elapsed since its last visit. Hence there is **no notion of memory** (estimating the expected idleness based on previous visits)

The name 2_sim_rl indicates => sim = simple estimation of expected idleness (AGENT model)
rl = using reinforcement learning (Q-learning algorithm to calculate the value function)

2. For a bot deciding the next node to visit

Set:

- True idleness of the node where the bot is present = 0
- Expected idleness of the node wrt the corresponding bot = 0

Decision Making: here, we chose ε=0.1

- With $(1 - \varepsilon)$probability, check all neighbours and visit the one with highest value of = $[expected\ idleness]\ x\ [value\_exp]$

- With ε probability, go to a random node

----------END----------

DRAWBACK:

- Even when |expect-true| = 1, reward = 0.